Autonomous AI Research Loops: A Self-Sustaining Architecture for Continuous Scientific Discovery

Hróar Þór Reynisson[1] and Claude (Anthropic)[2]
with the Sefirot Autonomous Research System

[1]Iðntæknistofnun (IceTec), Reykjavík, Iceland. ORCID: 0009-0003-6518-121X
[2]Anthropic, Claude 4.5 Sonnet family

Abstract

We present KETER, an architecture for autonomous AI research loops that enables continuous scientific investigation without human intervention. The system decomposes research capability into a domain-independent template—encompassing orchestration, quality control, knowledge management, and strategic self-assessment—and domain-specific plugins for constants, validation, mathematical tools, and publication formats. We validate the architecture through sustained operation in theoretical physics, where the system completed 100+ autonomous research cycles, accumulated 279 verified knowledge claims across a 4-tier evidence hierarchy, closed 10 research arcs, and produced a submission-ready manuscript—all without human intervention during operation. Component analysis reveals that 66% of the architecture is domain-independent, suggesting feasibility of instantiation for other research domains. Key innovations include: (1) reflection checkpoints with automated concept correlation for strategic self-assessment, (2) honest negative requirements enforcing scientific rigor, (3) succession protocols enabling indefinite cross-session operation, and (4) a quality tier system mapping evidence strength across domains. We compare KETER against existing autonomous research systems (AI Scientist, BabyAGI, AutoGPT) and identify sustained multi-cycle autonomy with quality control as the primary differentiator. We document five categories of honest negatives about the architecture itself, including the inherent circularity of self-referential validation.

1 Introduction

The automation of scientific research has progressed from individual tool assistance (citation managers, statistical packages) through semi-automated pipelines (AutoML [**?**], neural architecture search) to recent attempts at fully autonomous research paper generation [**?**]. However, existing systems operate in a *one-shot* paradigm: they generate a single paper from a single research question, with no mechanism for sustained investigation, knowledge accumulation, or strategic self-correction across multiple research cycles.

Real scientific research is not one-shot. It is an iterative process of hypothesis generation, testing, revision, and accumulation, often spanning years and thousands of individual experiments. A researcher's tenth paper in a field benefits from the knowledge, failed approaches, and refined intuitions accumulated across the first nine. No existing autonomous system captures this longitudinal dimension of research.

We present KETER, an architecture designed for *sustained autonomous research.* The system operates in continuous cycles—each comprising task dispatch, execution, result ingestion, and housekeeping—with periodic reflection checkpoints that assess strategic direction, identify productive and stuck research threads, and generate new research hypotheses from accumulated

data. The architecture enforces scientific rigor through mandatory honest negative documentation, a 4-tier evidence hierarchy, and quality gates that prevent overclaiming.

Our primary contribution is not generating papers (AI Scientist v2 does that) but sustaining a coherent research program across 100+ cycles with monotonically growing knowledge, strategic self-correction, and zero human intervention during operation.

Contributions.

1. A formal specification of a domain-independent autonomous research loop template with 10 clearly defined insertion points for domain-specific components (Section **??**).
2. Quantitative analysis from 100+ operational cycles in theoretical physics, demonstrating sustained knowledge accumulation (279 claims), quality distribution (23% BEDROCK, 55% PILLAR), and self-correction mechanisms (Section **??**).
3. Component analysis showing 66% of the architecture is domain-independent, with detailed mapping of shared versus domain-specific components across physics and computer science (Section **??**).
4. Honest assessment of five categories of limitations, including self-referential circularity, context window constraints, and domain coverage boundaries (Section **??**).

2 Related Work

Autonomous AI agents. AutoGPT [**?**] and BabyAGI [**?**] demonstrated autonomous task execution via LLM-driven loops with tool use. These systems decompose goals into subtasks and execute them iteratively but lack research-specific quality controls, knowledge persistence, or scientific rigor constraints. They operate without evidence tiers, honest negative requirements, or reflection checkpoints, making them unsuitable for sustained scientific investigation where accumulation and self-correction are essential.

AI-driven research. The AI Scientist [**?**] (Sakana AI, 2024) was the first system to automate the full research lifecycle from idea generation through paper writing and self-review. AI Scientist v2 [**?**] removed template dependencies and produced the first AI-authored workshop paper accepted through peer review. However, both versions operate in a one-shot paradigm: each run produces one paper from one idea, with no knowledge carried forward to subsequent runs. Independent evaluation revealed a 42% experiment failure rate and poor novelty assessment [**?**]. AIDE [**?**] focuses on AI-driven experiment design and optimization but does not produce research papers or maintain long-term knowledge.

Research methodology frameworks. SWE-Agent [**?**] and OpenHands [**?**] automate software engineering tasks (bug fixing, feature implementation) but target known problem resolution rather than open-ended scientific discovery. MLAgentBench [**?**] provides benchmarks for evaluating ML research agents but is an evaluation framework, not an autonomous research system.

Positioning. KETER differs from all the above in its focus on *sustained multi-cycle autonomy* with quality control. Where AI Scientist generates one paper per run, KETER maintains a research program across 100+ cycles with knowledge accumulation, reflection checkpoints, honest negative enforcement, and succession protocols for indefinite operation. Table **??** summarizes the key differences.

Table 1: Comparison of autonomous research systems.

Feature	KETER	AI Sci. v2	BabyAGI	AutoGPT	AIDE
Multi-cycle autonomy	100+	1	∞^*	∞^*	1
Quality tier system	4-tier	–	–	–	–
Honest negatives	≥3/unit	–	–	–	–
Reflection checkpoints	Every 5 cyc.	–	–	–	–
Knowledge accumulation	279+ claims	–	Task list	Memory	–
Self-correction	AMBER flags	Self-review	–	–	–
Concept correlation	Automated	–	–	–	–
Succession protocol	Cross-session	–	–	–	–
Arc organization	Open/close/grade	Linear	Flat	Goal tree	–

*Unbounded but undirected—no quality control or research structure.

3 Architecture

KETER decomposes autonomous research into a *domain-independent template* (66% of components) and *domain-specific plugins* (34%). We describe the template here; Section **??** covers instantiation.

3.1 Role Hierarchy

The system employs a 4-level hierarchy inspired by organizational delegation:

1. **CEO (Crown):** Sets research focus, manages priority stack, reads only briefings (≤100 lines). Never processes raw data.
2. **Director:** Dispatched per cycle. Scans task queue, dispatches workers, collects results, writes structured report.
3. **Workers:** Execute specific research tasks in parallel. Write raw output to persistent storage, return summaries (≤80 lines).
4. **Auditors:** Verify claims, run validation tests, enforce quality gates.

This hierarchy enforces *context isolation*: the CEO's context budget is protected by strict output limits at each level (the "Commando Pattern"), preventing the information overload that degrades LLM performance in long contexts.

3.2 Main Loop

Algorithm **??** shows the core loop. Each cycle comprises: (1) exit condition checks, (2) optional reflection checkpoint, (3) Director dispatch with parallel worker execution, (4) claim ingestion into persistent knowledge base, (5) version-controlled commit, (6) knowledge graph synchronization, and (7) process cleanup. The loop runs autonomously until a termination condition triggers or cycles are exhausted.

3.3 Quality Tier System

Every knowledge claim is assigned one of four evidence tiers:

- **BEDROCK:** Highest confidence. Independently reproducible with minimal uncertainty. *Domain-specific threshold.*
- **PILLAR:** Strong evidence. Reproducible with moderate uncertainty.
- **SAND:** Preliminary. Partial evidence or limited evaluation.
- **MUD:** Speculative. Insufficient evidence for any confidence.

The tier definitions are domain-specific insertion points. In physics, BEDROCK means <1 ppm agreement with authoritative measurement. In CS, it might mean exceeding state-of-the-art with

Algorithm 1 Keter Autonomous Research Loop

```
Require: mode ∈ {fixed, auto}, max_cycles ∈ ℤ⁺
cycle ← 0
while cycle < max_cycles do
    cycle ← cycle + 1
    if 3-strike rule triggered then
        exit("context_death_spiral")
    end if
    if cycle mod N = 0 and max_cycles ≥ N then
        report ← ReflectionCheckpoint(cycle)
        if report.escalation_needed then
            exit("reflection_escalation")
        end if
    end if
    director_report ← DispatchDirector(cycle, priority_stack)
    if director_report contains critical discovery then
        exit("critical_discovery")
    end if
    IngestClaims(director_report.claims, knowledge_base)
    Commit("cycle-{cycle}: {summary}")
    SyncKnowledgeGraph()
    ProcessCleanup()
end while
WriteSessionReport()
```

$p < 0.001$ on a standard benchmark. The *structure* (4 ordered tiers with promotion/demotion rules) is universal.

3.4 Honest Negative Requirements

Every investigation unit must document ≥ 3 *honest negatives*: falsified hypotheses, boundary conditions, limitations, or failed approaches. This is not a suggestion but an architectural constraint—investigation outputs missing honest negatives are rejected at the quality gate.

This requirement prevents the well-documented tendency of LLMs to overclaim results and suppress negative findings [**?**]. In 100+ operational cycles, the system documented 500+ honest negatives, including the demotion of its own axioms (discovering circular reasoning in foundational claims).

3.5 Reflection Checkpoints

Every Nth cycle (default $N = 5$), the main loop pauses for strategic self-assessment. A 4-agent team executes in sequence:

1. **State Auditor:** Inventories current knowledge (claim counts by tier, test pass rates, active arcs).
2. **Concept Correlator:** Extracts keywords from accumulated claims, builds co-occurrence matrices, identifies unexpected correlations, and auto-generates research hypotheses ("seeds").
3. **Strategic Evaluator:** Assesses productive vs. stuck threads, applies quality gates (e.g., GPA recovery for underperforming arcs), and flags escalation conditions.
4. **Report Compiler:** Produces three-tier reporting: Executive (full archive), Brief ($\leq$50 lines for CEO), and Director Index (machine-parseable task list).

Table 2: Domain-specific insertion points in the KETER template.

Insertion Point	Physics Instance	CS Instance
Constants module	CODATA physical constants	Benchmark parameters
Validation tests	1716 physics tests	Benchmark suite
Crown jewels	m_p/m_e at 0.015 ppb	Sustained autonomy metrics
Tag taxonomy	e8, qed, mass_ratio	deep_learning, NLP
Quality gates	Alpha decision gate	Significance threshold
Math infrastructure	E8, Clifford algebras	ML frameworks, datasets
Publication format	RevTeX (APS journals)	NeurIPS style file
Data sources	arXiv hep-th, INSPIRE	Semantic Scholar
Experiment runner	Simulation scripts	Benchmark harness
Concept type defs.	Physical law / formula	Algorithm / technique

The Concept Correlator is a key innovation: by analyzing cross-claim keyword co-occurrence, it identifies research directions the system would not have considered through linear task execution alone.

3.6 Knowledge Persistence

Claims are stored in a persistent knowledge base with full dependency tracking:

- **Graph structure:** Claims as nodes; edges typed as `depends_on`, `supports`, `derives`, `contradicts`, `refines`.
- **Concept types:** `prime` (foundational), `edge` (connecting), `derived` (dependent).
- **Gap detection:** Automated identification of missing dependencies and orphan claims.
- **Hierarchical context:** Nested graph mapping arcs $\supset$ phases $\supset$ experiments $\supset$ claims.

3.7 Domain-Specific Insertion Points

Table **??** lists the 10 insertion points where domain-specific content plugs into the universal template.

4 Case Study: Theoretical Physics

We validate the architecture through sustained operation in theoretical physics (the QHOTS project: Qameah Harmonic Oscillator Theory of Symmetry), where the system investigated geometric derivations of fundamental physical constants.

4.1 Operational Metrics

- **Cycles:** 100+ autonomous cycles across 7 sessions
- **Phases:** 400+ research phases completed
- **Experiments:** 963 simulations executed
- **Knowledge:** 279 claims accumulated (64 BEDROCK, 153 PILLAR, 54 SAND, 8 MUD)
- **Arcs:** 10 research arcs closed, 7 graded at GPA ≥ 3.0
- **Tests:** 1716 validation tests, 294 in campaign framework
- **Honest negatives:** 500+ documented across all phases
- **Self-corrections:** Multiple axiom demotions (phi circularity, 28 continued fraction origin)
- **Output:** 1276-line submission-ready manuscript (62 revisions)
- **Infrastructure recoveries:** Survived segfaults, orphaned commits, context overflows—all recovered autonomously

4.2 Quality Distribution

The tier distribution of accumulated claims follows a pyramid structure (Figure **??**): 23% BEDROCK (64), 55% PILLAR (153), 19% SAND (54), 3% MUD (8). This distribution is *not* by design—the system does not target tier ratios. Rather, it reflects the natural difficulty of achieving high-confidence results: most findings are PILLAR-tier (strong but not definitive), with a smaller proportion reaching BEDROCK status.

The 3% MUD rate indicates the system rarely generates speculative claims without evidence, suggesting the honest negative requirements and quality gates are effective filters.

4.3 Self-Correction in Action

A critical test of autonomous research is whether the system can identify and correct its own errors. Three notable self-corrections occurred:

1. **Phi circularity demotion:** The system discovered that its use of the golden ratio ϕ in certain derivations was circular (the conclusion assumed the premise). It autonomously demoted the affected claims from PILLAR to SAND and documented the circularity.
2. **28 continued fraction origin:** A claim about the nuclear magic number 28 arising from a continued fraction was found to have an alternative, simpler explanation. The system documented both explanations and flagged the original as "selection bias."
3. **AMBER flag system:** Editorial drift (gradual inflation of claims over many cycles) was detected by the reflection checkpoints and corrected via mandatory "pivot" cycles that re-evaluated accumulated evidence.

5 Results: Architecture Analysis

5.1 Component Classification

We audited all 68 identifiable components of the operational system and classified each as UNIVERSAL (domain-independent) or DOMAIN-SPECIFIC:

- **Universal:** 45 components (66.2%)—including role hierarchy, main loop, reflection system, quality tiers, honest negatives, context management patterns, knowledge base schema, skill forge, and anti-patterns.
- **Domain-specific:** 23 components (33.8%)—including physical constants, validation test content, mathematical infrastructure, crown jewel definitions, and publication format.

The universal components form a self-consistent framework. A reader unfamiliar with the source domain could instantiate the template for a new domain by filling 10 insertion points (Table **??**).

5.2 Cross-Domain Mapping

We analyzed the feasibility of mapping 8 shared components from physics to CS (Table **??**). Three components map cleanly (honest negatives, concept correlation, reflection checkpoints). Two have minor mismatches (publication pipeline needs page-limit awareness; arc organization needs deadline awareness). Three have significant mismatches requiring adaptation:

5.3 Sustainability Metrics

Over 100+ cycles, we observe:

- **Knowledge growth:** Approximately linear (2–3 claims/cycle), with no sign of saturation.
- **Quality maintenance:** BEDROCK proportion remained stable at 20–25% throughout, not declining with volume.

Table 3: Cross-domain component mapping quality.

Component	Mapping Quality	Key Mismatch
Honest negatives	Clean	None
Concept correlation	Clean	None
Reflection checkpoints	Clean	Deadline awareness needed
Publication pipeline	Minor	Page limits require budget pass
Arc organization	Minor	Conference deadlines force early closure
Quality tiers	Significant	No single authoritative reference (no "CODATA")
Knowledge base	Significant	CS results expire (staleness mechanism needed)
Validation framework	Significant	Multi-objective evaluation (accuracy + latency + fairness)

- **Self-correction rate:** Reflection checkpoints caught 3 major errors (circularity, selection bias, editorial drift) that would have compounded without intervention.
- **Infrastructure resilience:** The system survived 5 infrastructure failures (segfaults, orphaned commits, context overflows) and recovered autonomously in all cases via checkpoint-based restart.

6 Analysis: What Works and What Fails

What works.

1. **Concept correlation discovers non-obvious connections.** The automated keyword co-occurrence analysis identified research directions (e.g., connecting E8 geometry to neutron mass formulas) that sequential task execution would not have explored.
2. **Honest negatives prevent overclaiming.** The mandatory $\geq$3 negatives per unit forced the system to document boundaries and limitations that improved the quality of surviving claims.
3. **Tier system enables graduated confidence.** Rather than binary accept/reject, the 4-tier system allows preliminary findings (SAND) to coexist with verified results (BEDROCK), supporting exploratory research without conflating evidence levels.

What fails.

1. **Editorial drift.** Over many cycles, claims gradually inflated in confidence without corresponding evidence improvement. The AMBER flag system in reflection checkpoints partially mitigates but does not eliminate this tendency.
2. **Context window limits constrain cycle depth.** Each cycle must fit within the LLM's context window. Complex investigations requiring deep multi-step reasoning across many files are truncated, leading to shallow analysis in some phases.
3. **Diminishing returns in mature arcs.** After 5+ phases in a single arc, the discovery rate (new BEDROCK claims per phase) declined. The system mitigated this by closing arcs and opening new ones, but the optimal arc length is not yet understood.

What self-corrects. The reflection checkpoint system enables self-correction at the strategic level: identifying stuck threads, triggering arc closure, generating new research directions via concept correlation, and flagging GPA degradation for recovery action. This operates at a 5-cycle cadence, which empirically balances assessment frequency against overhead cost.

7 Limitations

We document five categories of limitations as honest negatives about the architecture itself:

1. Context window constraints. The fundamental bottleneck is the LLM's context window. Despite context management patterns (CEO Pattern, Commando Pattern, disk-first communication), complex investigations that require simultaneous awareness of many files or long derivation chains are limited by the window size. The succession protocol (cross-session continuity via checkpoint files) partially addresses this but loses in-context reasoning state.

2. Self-referential circularity. Any system that evaluates its own research output faces inherent circularity. The quality tier assignments, GPA grades, and reflection assessments are all produced by the same LLM that generated the research. External validation (human expert review, independent reproduction, venue peer review) is the only true escape from this circularity. We do not claim the system is "self-validating"—we claim it is *self-correcting within the limits of its own capability*, which is a weaker but still useful property.

3. Domain coverage boundary. The architecture has been validated in one domain (theoretical physics). While component analysis suggests 66% domain-independence, actual instantiation in a second domain has not been completed. Claims about generality are projections, not demonstrations. The three significant component mismatches identified in Section **??** (staleness, multi-objective evaluation, no single authoritative reference) may reveal deeper obstacles during actual CS instantiation.

4. Model capability ceiling. The system's research quality is bounded by the underlying LLM's capabilities. It cannot discover mathematics it cannot derive, cannot reason about physics it doesn't understand, and cannot evaluate literature it hasn't been trained on. The architecture amplifies and systematizes the model's capabilities but does not transcend them.

5. Reproducibility challenges. While the system commits every cycle's outputs to version control (enabling full audit trails), the stochastic nature of LLM generation means that rerunning the same cycle does not produce identical outputs. Reproducibility is at the *methodology* level (same architecture, same constraints, same quality thresholds) rather than the *output* level (same generated text).

8 Discussion

Implications for autonomous AI research. The KETER architecture suggests that sustained autonomous research is feasible when the system includes quality control, self-assessment, and knowledge accumulation—features absent from most existing autonomous agents. The key insight is that *research methodology* (how to do research) is largely domain-independent, while *research content* (what to research) is domain-specific. This separation enables a template-based approach where the methodology framework is reused and only the domain content is replaced.

The quality tier innovation. The 4-tier evidence hierarchy addresses a fundamental challenge in LLM-driven research: the tendency to present all outputs with equal confidence. By forcing explicit tier assignment with defined thresholds, the system creates graduated confidence that more closely mirrors human scientific judgment.

Honest negatives as architectural constraint. Making honest negatives mandatory (rather than optional) transforms them from a reporting afterthought into a structural feature. The 500+ honest negatives accumulated across 100+ cycles represent a valuable dataset about the boundaries of the research approach—arguably more informative than the positive results alone.

Toward indefinite operation. The succession protocol—where the system writes a "startup plan" for future sessions when context resources are depleted—enables theoretically indefinite operation. Each session reads the previous session's checkpoint and continues from where it left off. In practice, this has been demonstrated across 7 sessions with no loss of strategic direction.

9 Self-Referential Validation

This section acknowledges the strongest and most problematic aspect of this work: the system that produced the research described in this paper is the same system described by this paper. This is simultaneously the strongest possible proof of concept (the architecture works well enough to produce a paper about itself) and the weakest possible validation (the system grades its own output).

We address this tension directly:

1. **The architecture existed before this paper.** KETER was designed and deployed for physics research, not for self-description. The template extraction and CS domain analysis are post-hoc analysis of an operational system, not a purpose-built demonstration.
2. **External validation is required.** We explicitly do not claim this paper constitutes proof of the architecture's quality. Venue peer review, independent reproduction by other research groups, and instantiation in a third domain (neither physics nor CS meta-analysis) would constitute meaningful external validation.
3. **Failure would be informative.** If this paper is rejected by peer review or found to be methodologically weak, that result documents where autonomous AI research breaks down—which is itself a contribution to the field. The architecture's honest negative requirements apply to itself.

10 Conclusion

We presented KETER, a domain-independent architecture for sustained autonomous AI research. Through 100+ operational cycles in theoretical physics, the system demonstrated knowledge accumulation (279 claims), quality control (4-tier evidence hierarchy with 23% BEDROCK rate), self-correction (axiom demotion, editorial drift detection), and infrastructure resilience (5 failure recoveries). Component analysis reveals 66% domain-independence with 10 clearly defined insertion points for domain-specific content.

The primary contribution is shifting the autonomous research paradigm from one-shot paper generation to sustained, self-correcting investigation with quality control and knowledge accumulation. The primary limitation is that validation in a single domain, while extensive, does not prove generality.

Future work includes: (1) instantiation for CS research to test cross-domain generality, (2) formal analysis of optimal reflection checkpoint frequency, (3) multi-agent negotiation for conflicting claims, and (4) integration with external validation services for non-circular quality assessment.

References

[1] H. He, G. Hutterer, and Q. Yao, "AutoML: A survey of the state-of-the-art," *Knowledge-Based Systems*, vol. 212, 106622, 2021.

[2] C. Lu, C. Lu, R. T. Lange, J. Foerster, J. Clune, and D. Ha, "The AI Scientist: Towards fully automated open-ended scientific discovery," *arXiv:2408.06292*, 2024.

[3] C. Lu *et al.*, "The AI Scientist-v2: Workshop-level automated scientific discovery via agentic tree search," *Sakana AI Technical Report*, 2025.

[4] M. Y. Kan *et al.*, "Evaluating Sakana's AI Scientist for autonomous research," *arXiv:2502.14297*, 2025.

[5] T. B. Richards, "Auto-GPT: An autonomous GPT-4 experiment," *GitHub Repository*, 2023.

[6] Y. Nakajima, "BabyAGI," *GitHub Repository*, 2023.

[7] Weco AI, "AIDE: AI-driven experiment design," *GitHub Repository*, 2024.

[8] C. E. Jimenez *et al.*, "SWE-agent: Agent-computer interfaces enable automated software engineering," *arXiv:2405.15793*, 2024.

[9] X. Wang *et al.*, "OpenHands: An open platform for AI software developers as generalist agents," *arXiv:2407.16741*, 2024.

[10] Q. Huang *et al.*, "MLAgentBench: Evaluating language agents on machine learning experimentation," *ICML*, 2024.

[11] Y. Huang *et al.*, "A survey on hallucination in large language models," *ACM Computing Surveys*, vol. 57, no. 2, 2024.